VerifAI: all datasets and where to find them

Green rows are organic (public, written by outsiders), purple rows are synthetic (generated by templates or an LLM), white rows are handwritten by the team. Licence notes are what I expect; confirm each on its dataset card before reusing the data.

#	Dataset / model	Type	Link or location	Rows	How VerifAI uses it	Licence / access
1	**XSTest**	Organic	https://huggingface.co/datasets/walledai/XSTest	450 (250 safe look-alikes, 200 unsafe)	Half trains the policy model; other half is held-out test	CC-BY-4.0 (check the card); needs Hugging Face login
2	**JailbreakBench (JBB-Behaviors)**	Organic	https://huggingface.co/datasets/JailbreakBench/JBB-Behaviors	100 harmful + 100 benign	Half trains, half held-out test	MIT (check the card)
3	**AgentHarm**	Organic	https://huggingface.co/datasets/ai-safety-institute/AgentHarm	208 harmful agent tasks	Half trains, half held-out test	Custom licence, gated; check the card
4	**Banking77**	Organic	https://huggingface.co/datasets/PolyAI/banking77	1,000 train sample, 300 test sample	Harmless real banking queries (train negatives and held-out test)	CC-BY-4.0 (check the card)
5	**WildGuard (safety model)**	Organic model	https://huggingface.co/allenai/wildguard	7B model	Default safety judge (port 8000)	AI2 licence; gated, needs Hugging Face login
6	**Llama Guard 3 1B (MLX 4-bit)**	Organic model	https://huggingface.co/mlx-community/Llama-Guard-3-1B-chat-mlx-4bit	1B model	Optional smaller safety judge (port 8002)	Llama 3.2 community licence
7	**Llama 3.1 8B Instruct (MLX 4-bit)**	Organic model	https://huggingface.co/mlx-community/Meta-Llama-3.1-8B-Instruct-4bit	8B model	Optional local agent (port 8001)	Llama 3.1 community licence
8	**policy_intent.jsonl**	Synthetic	training/data/policy_intent.jsonl (made by training/generate_policy_data.py)	6,965	Trains the policy model (templates and harmless chit-chat)	Own data
9	**policy_llm.jsonl**	Synthetic	training/data/policy_llm.jsonl (made by training/generate_policy_llm.py)	3,118	Trains the policy model (LLM-written rewordings)	Own data; written by Groq gpt-oss-20b
10	**generated.jsonl**	Synthetic	training/data/generated.jsonl (made by training/generate_data.py)	3,489	Trains the impact model	Own data; written by Groq gpt-oss-20b
11	**generated_focus.jsonl**	Synthetic	training/data/generated_focus.jsonl (made by training/generate_focus.py)	2,831	Trains the impact model (weak spots)	Own data
12	**eval_verdicts, _2, _3**	Synthetic	tests/data/eval_verdicts*.jsonl (made by training/generate_eval_sets.py)	3 x 180	Tests the whole system (allow / review / block)	Own data
13	**eval_tools, _2, _3**	Synthetic	tests/data/eval_tools*.jsonl (made by training/generate_eval_sets.py)	170 + 170 + 168	Tests the agent's tool choice	Own data
14	**Groq gpt-oss-20b (writer of the above)**	Synthetic source	https://console.groq.com/docs/model/openai/gpt-oss-20b	n/a	LLM that wrote the synthetic data and runs as the agent	Groq API key required
15	**handwritten_test(2).jsonl**	Handwritten	training/data/handwritten_test.jsonl and handwritten_test2.jsonl	83 + 60	Tests the impact model	Own data
16	**hard_policy_cases.json**	Handwritten	tests/hard_policy_cases.json	64 + 65	Tests the pattern rules	Own data

#	Dataset / model	Type	Link or location	Rows	How VerifAI uses it	Licence / access
17	**hard_policy_paraphrases.json**	Handwritten	tests/hard_policy_paraphrases.json	63 + 20	Tests the policy model	Own data

Synthetic and handwritten data have no web link: they are files inside the project folder, made by the generator scripts named in the table. Not used: OmniCompliance (discussed but not trained on).